

Inverse Problems in Photonics

Machine Learning Accelerated Solutions of Inverse Problems in Photonics

Gustav Fredrikson
Johan Rensfeldt
Fredrik Gillgren

Francesco Ferranti & Prashant Singh
Uppsala University

January 8, 2024

Table of Contents

- 1 Introduction
- 2 Previous Research
- 3 Theory & Method
- 4 Results & Discussion
- 5 Conclusions

Introduction

Introduction

Research Focus

- Addressing inverse problems in photonics with advanced machine learning techniques.
- Exploring Mixture Density Network (MDN), Invertible Neural Network (INN), and Dirichlet Process Gaussian Mixture Model (DPGMM).
- Tackling the one-to-many mapping problem in nanophotonic structures.




Figure: Schematic representation of our methodological approach, showing the stages from data preparation to model evaluation and testing.

Introduction - Dataset Overview

Data Composition

- The dataset is sourced from CSV files.
- Input features X and target responses Y .
- Dimensions of X : 3847×5 (samples \times features).
- Dimensions of Y : 3847×101 (samples \times features).

Data Processing and Normalization

- Dataset split into training, validation, and testing subsets.
- Normalization: Both X and Y values are normalized between 0 and 1.

Previous Research

Our contribution



Theory & Method

Theory - MDN - Overview

- **Mixture density network (MDN):** A Mixture Density Network is a type of neural network that outputs the parameters of a mixture model, usually Gaussian, to model complex and multimodal data distributions.

Mathematical Formulation

$$p(\mathbf{t} \mid \mathbf{x}) = \sum_{i=1}^m \alpha_i(\mathbf{x}) \phi_i(\mathbf{t} \mid \mathbf{x}) \quad (1)$$

$$\phi_i(\mathbf{t} \mid \mathbf{x}) = \frac{1}{(2\pi)^{\frac{c}{2}} \sigma_i(\mathbf{x})^c} \exp \left\{ -\frac{\|\mathbf{t} - \mu_i(\mathbf{x})\|^2}{2\sigma_i(\mathbf{x})^2} \right\} \quad (2)$$

$$\sum_{i=1}^m \phi_i(\mathbf{x}) = 1, \alpha_i = \frac{\exp(z_i^\alpha)}{\sum_{i=1}^M \exp(z_i^\alpha)}, \sigma_i = \exp(z_i^\sigma), \mu_{ik} = z_{ik}^\mu \quad (3)$$

Theory - MDN - Model structure

- 1 Neural network model learns parameters: means, variances, and mixture coefficients for each component.
- 2 Mixture model uses the parameters output by the neural network to define a set of Gaussian distributions.
- 3 Making predictions?

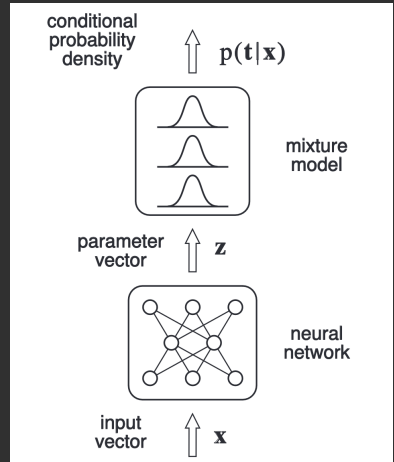
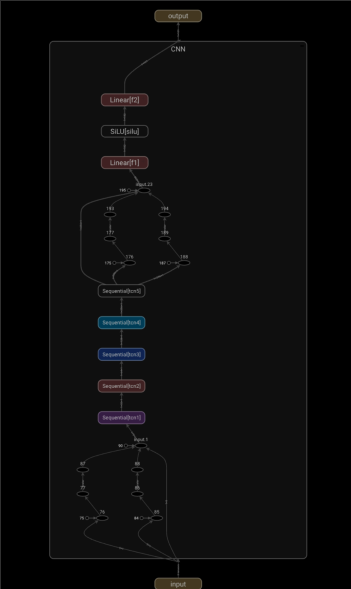



Figure: Model overview

Theory & Method

Theory - MDN - Making predictions

- # Theory - MDN - Making predictions
- How do you choose what the most probable value of x given t ?
 - 1 The most likely value of x given input vector t is the maximum of the conditional density $p(t | x)$.
 - 2 A good approximation is $\max(\frac{\alpha_i(x^c)}{\sigma_i(x^c)})$.
 - 3 Or training another neural network in choosing the best approximate value.
- 
- The diagram illustrates a deep neural network architecture. At the bottom, an 'Input' node feeds into a series of layers. The first layer is a 'Sequential' block (purple). This is followed by three 'Sequential' blocks (red, blue, and blue). Then, there is a 'Sequential' block (red) that branches into two parallel paths. Each path contains a 'ReLU' activation function followed by a 'Linear' layer. The outputs of these two paths are combined. This is followed by another 'Sequential' block (blue), then a 'ReLU' activation function, and finally a 'Linear' layer (red). The output of this final 'Linear' layer is fed into a 'CNN' block, which then produces the final 'Output'.



Method - MDN - Overview

- 1 Create and train several different MDNs (MDN1-6).
- 2 Optimize models
- 3 Evaluate Model accuracy on all of the models individually.
- 4 Combine the models.

```

1 class MDN3(nn.Module):
2     def __init__(self, in_features=10, out_features=5, num_gaussians=9):
3         super(MDN3, self).__init__()
4
5         self.in_features, self.out_features, self.num_gaussians = in_features, out_features, num_gaussians
6
7         # Different architecture for shared layers
8         self.shared_layers2 = nn.Sequential(
9             nn.Conv1d(in_features, 256, kernel_size=3, stride=1, padding=1),
10             nn.BatchNorm1d(256),
11             nn.Dropout(0.2),
12             nn.ReLU(),
13             nn.Conv1d(256, 256, kernel_size=3, stride=1, padding=1),
14             nn.BatchNorm1d(256),
15             nn.Dropout(0.2),
16             nn.ReLU(),
17         )
18
19         self.shared_layers1 = nn.Sequential(
20             nn.Linear(256, 256),
21             nn.BatchNorm1d(256),
22             nn.ReLU(),
23             nn.Dropout(0.2),
24             nn.Linear(256, 128),
25             nn.BatchNorm1d(128),
26             nn.ReLU(),
27             nn.Dropout(0.2),
28             nn.Linear(128, 128),
29             nn.BatchNorm1d(128),
30             nn.ReLU(),
31             nn.Dropout(0.2),
32             nn.Linear(128, 64),
33             nn.ReLU(),
34         )
35
36         # Different activation function for pi and a different structure for sigma and mu
37         self.pi = nn.Sequential(
38             nn.Linear(64, num_gaussians),
39             nn.Softmax(dim=1)
40         )
41         self.sigma = nn.Linear(64, out_features * num_gaussians)
42         self.mu = nn.Linear(64, out_features * num_gaussians)
43
44         def forward(self, y):
45             y = y.reshape(y.shape[0], y.shape[1], 1)
46             y = self.shared_layers2(y)
47             y = y.reshape(y.shape[0], y.shape[1])
48             y = self.shared_layers1(y)
49             pi = self.pi(y)
50             sigma = torch.exp(self.sigma(y))
51             sigma = sigma.view(-1, self.num_gaussians, self.out_features)
52             mu = self.mu(y)
53             mu = mu.view(-1, self.num_gaussians, self.out_features)
54             return pi, sigma, mu
55
56         def init_weights(self, m):

```

Method - MDN - Training models

- Using ADAMS optimizer.
- Negative log likelihood loss function.
- Using gradient clipping.
- Using StepLR scheduler.
- 500 epochs

```

1 def train_mdn(t_loader, v_loader, n_epochs, early_stop, model):
2     if isinstance(model, MDN4) or isinstance(model, MDN6):
3         optimizer = torch.optim.Adam(model.parameters(), lr=1e-3, weight_decay=1e-4, eps=1e-8, betas=(0.9, 0.999))
4         scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=100, gamma=0.5)
5     else:
6         optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)
7     early_stop_counter = 0
8     best_model_wts = copy.deepcopy(model.state_dict())
9     best_loss, best_train = 1000.0, 0
10    history = dict(train=[], val=[])
11    model.apply(model.init_weights)
12    for epoch in range(1, n_epochs + 1):
13        # Training
14        model = model.train()
15        train_losses = []
16        for batch_idx, (y, x) in enumerate(t_loader):
17            optimizer.zero_grad()
18            pi_variable, sigma_variable, mu_variable = model(y)
19            loss = mdn_loss_fn(pi_variable, sigma_variable, mu_variable, x)
20            loss.backward()
21            optimizer.step()
22            train_losses.append(loss.item())
23        if isinstance(model, MDN4) or isinstance(model, MDN6):
24            old_lr = scheduler.get_last_lr()[0] # Get the last learning rate
25            scheduler.step()
26            new_lr = scheduler.get_lr()[0] # Get the current learning rate
27            if new_lr != old_lr:
28                print(f"Epoch {epoch}: Learning rate changed from {old_lr} to {new_lr}")
29        if isinstance(model, MDN4) or isinstance(model, MDN6):
30            # gradient clipping triggered when gradient explodes
31            for p in model.parameters():
32                if p.grad is not None:
33                    if torch.any(torch.isnan(p.grad)):
34                        print(f"Epoch {epoch}: Gradient exploded")
35                        nn.utils.clip_grad_norm_(model.parameters(), 1.0)
36

```

Figure: Model overview

Method - MDN - Optimizing models

- Focusing on dropout rate and the number of gaussians.
- Baesian gridsearch.
- Ax and BoTorch.
- 40 trails and 300 epochs.

```

1 def grid_search_bayesian_parallel(models, num_trials=30, num_epochs=250, early_stop=5, batch_size=5):
2     all_trial_mappings = {}
3     best_model_params = {}
4     for model in models:
5         model_name = model.__name__
6         print(f'Now training {model_name}')
7         ax_client = AxClient()
8         ax_client.create_experiment(
9             name=f"mdn_experiment_{model_name}",
10            parameters=[
11                ("num_gaussians", "type": "range", "bounds": [6, 14], "value_type": "int"),
12                ("dropout", "type": "range", "bounds": [0.1, 0.8], "value_type": "float"),
13            ],
14            objective_name="validation_loss",
15            minimize=True
16        )
17        trial_model_mapping = {}
18        for batch in tqdm(range(0, num_trials, batch_size), desc=f"Optimizing {model_name}"):
19            trials_data = []
20            for _ in range(batch_size):
21                parameters, trial_index = ax_client.get_next_trial()
22                trials_data.append((parameters, trial_index))
23                unique_trial_index = f"{model_name}_{batch + 1}_{trial_index}"
24                trial_model_mapping[unique_trial_index] = model
25                # Training models in parallel for the current batch
26                validation_losses = train_model_in_parallel(model, [data[0] for data in trials_data],
27                                                            mdn_train_loader, mdn_val_loader, num_epochs,
28                                                            early_stop)
29                # Completing each trial in the batch
30                for (validation_loss, (_, trial_index)) in zip(validation_losses, trials_data):
31                    ax_client.complete_trial(trial_index=trial_index, raw_data=(validation_loss, 0.0))
32            best_parameters, values = ax_client.get_best_parameters()
33            best_model_params[model_name] = best_parameters
34            all_trial_mappings[model_name] = trial_model_mapping
35
36    return best_model_params, all_trial_mappings

```

Figure: Model overview

Method - MDN - Models

■ Summary of the models

Model	Num. Gaussians	Special Features
MDN1	12	Linear layers with SiLU, Tanh
MDN2	8	Conv1d, BatchNorm1d, ReLU
MDN3	12	Conv1d, ELU, Linear
MDN4	13	LSTM layers, ReLU, Linear
MDN5	10	Linear layers with LeakyReLU, PReLU
MDN6	12	Conv1d, MaxPool1d, Relu

Table: Overview of MDN Models

Method - MDN - Combining models

- Average mu, pi and sigmas or average model outputs.

```

1 # Load models
2 models = [torch.load(f'MDN{i}.pth') for i in range(1, 7)]
3 # Load the forward model
4 mod = torch.load('models/forward_models/cnn.pth')
5 # Store final predictions from each model
6 all_final_predictions = []
7 # Generate and process predictions for each model
8 for model in models:
9     dS, dI81 = [], []
10    for batch_idx, (y, x) in enumerate(mdn_test_loader):
11        for i in y:
12            pi_variable, sigma_variable, mu_variable = model[i].unsqueeze(0)
13            j = list(pi_variable.detach().numpy()[0]).index(max(list(pi_variable.detach().numpy()[0])))
14            dS.append(list(mu_variable.detach().numpy()[0][j]))
15            pre = mod(torch.from_numpy(mu_variable.detach().numpy()[0][j]).unsqueeze(0)).squeeze().tolist()
16            dI81.append(pre)
17        all_final_predictions.append(dI81)
18    # Convert to numpy array for averaging
19    all_final_predictions_np = np.array(all_final_predictions)
20    # Average the final predictions
21    dI81 = np.mean(all_final_predictions_np, axis=0)

```

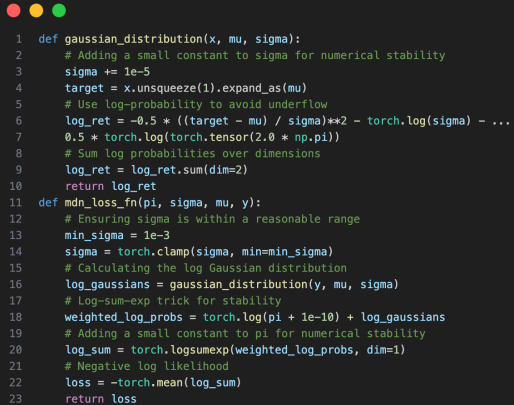
```

1 def ensemble_predict(models, input_data):
2     batch_pi, batch_sigma, batch_mu = [], [], []
3     for model in models:
4         pi, sigma, mu = model(input_data)
5         batch_pi.append(pi)
6         batch_sigma.append(sigma)
7         batch_mu.append(mu)
8     # Average across models
9     avg_pi = torch.mean(torch.stack(batch_pi), dim=0)
10    avg_sigma = torch.mean(torch.stack(batch_sigma), dim=0)
11    avg_mu = torch.mean(torch.stack(batch_mu), dim=0)
12    return avg_pi, avg_sigma, avg_mu
13 p, s, mu, o = [], [], [], []
14 for batch_idx, (y, x) in enumerate(mdn_test_loader):
15     for i in y:
16         pi_variable, sigma_variable, mu_variable = ensemble_predict(models, i.unsqueeze(0))
17         p.append(pi_variable.detach().numpy()[0])
18         s.append(sigma_variable.detach().numpy()[0])
19         m.append(mu_variable.detach().numpy()[0])
20 dS, dI81 = [], []
21 mod = torch.load('models/forward_models/cnn.pth')
22 for i in range(len(m)):
23     j = list(p[i]).index(max(list(p[i])))
24     dS.append(list(m[i][j]))
25     pre = mod(torch.from_numpy(m[i][j]).unsqueeze(0)).squeeze().tolist()
26     dI81.append(pre)

```


Method - MDN - Numerical stability

- 1 Using the Log-Sum-Exp Trick.
- 2 Adding a small constant to sigma and pi.
- 3 Avoiding numerical underflow by computing the log of the Gaussian probability rather than the probability itself.
- 4 Clamping sigma to a smallest value of 1e-3



```

1  def gaussian_distribution(x, mu, sigma):
2      # Adding a small constant to sigma for numerical stability
3      sigma += 1e-5
4      target = x.unsqueeze(1).expand_as(mu)
5      # Use log-probability to avoid underflow
6      log_ret = -0.5 * ((target - mu) / sigma)**2 - torch.log(sigma) - ...
7      0.5 * torch.log(torch.tensor(2.0 * np.pi))
8      # Sum log probabilities over dimensions
9      log_ret = log_ret.sum(dim=2)
10     return log_ret
11
12 def mdn_loss_fn(pi, sigma, mu, y):
13     # Ensuring sigma is within a reasonable range
14     min_sigma = 1e-3
15     sigma = torch.clamp(sigma, min=min_sigma)
16     # Calculating the log Gaussian distribution
17     log_gaussians = gaussian_distribution(y, mu, sigma)
18     # Log-sum-exp trick for stability
19     weighted_log_probs = torch.log(pi + 1e-10) + log_gaussians
20     # Adding a small constant to pi for numerical stability
21     log_sum = torch.logsumexp(weighted_log_probs, dim=1)
22     # Negative log likelihood
23     loss = -torch.mean(log_sum)
24     return loss
  
```

Figure: Model overview

DPGMM - Theory - Dirichlet Process

- **Dirichlet Process (DP):** A Bayesian nonparametric approach for modeling infinite-dimensional probability spaces.
- **Key Elements:** Base Distribution G_0 and Concentration Parameter α .

Mathematical Formulation

DP uses a stick-breaking process for its discrete nature:

$$G = \sum_{k=1}^{\infty} \beta_k \delta_{\theta_k}, \quad \beta_k = \nu_k \prod_{l=1}^{k-1} (1 - \nu_l), \quad \nu_k \sim \text{Beta}(1, \alpha). \quad (4)$$

DPGMM - Theory - GMM & DPGMM Combination

- **Gaussian Mixture Model (GMM):** A probabilistic model using a combination of Gaussian distributions.
- **Components:** Each Gaussian is characterized by a mean vector (μ_k) and covariance matrix (Σ_k).
- **DPGMM:** Merges the Dirichlet Process and Gaussian Mixture Model, leveraging nonparametric priors for the mixing proportions.

Mathematical Representation of DPGMM

DPGMM combines infinite Gaussian components with DP-derived mixing coefficients:

$$p(x) = \sum_{k=1}^{\infty} \pi_k \mathcal{N}(x | \mu_k, \Sigma_k), \quad (5)$$

where π_k are mixing coefficients from DP, and $\mathcal{N}(x | \mu_k, \Sigma_k)$ are Gaussian components.

DPGMM - Method - Class Structure and Functionality

■ Initialization:

- Initializes model with parameters: number of components, covariance type, weight concentration prior type, etc.
- Parameters define model behavior and capabilities.

■ Predict Method:

- Central to VBGMR class.
- Takes dataset and indices of input/output variables.
- Employs trained Gaussian mixture model for output estimation.

■ Model Training:

- 'fit' method trains VBGMR on a dataset.
- Optimizes mixture model parameters using variational Bayesian approach.

DPGMM- Method - Training and Evaluation Process

- Detailed process to train and evaluate the DPGMM.

Step	Description
1. Data Standardization	Scale invariance for input/output data.
2. Model Fitting	Training VBGMR on the dataset.
3. Analysis	Conducting forward and inverse analyses.
4. Performance Metrics	Evaluating with MSE and different losses.

Table: Training and Evaluation Steps for DPGMM

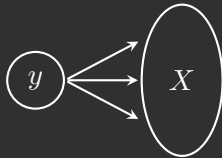
DPGMM - Method - Implementation Challenges

Implementation Challenges

- Balancing model complexity with performance for diverse data patterns.
- Achieving computational efficiency using Python and standard libraries.
- Potential enhancement with neural network-based training.


Theory - INN - Introduction

Original inverse problem



One-to-many mapping

Augmented inverse problem



Latent variable

Bijective mapping

The original problem is often ill-posed due to one-to-many mapping. An augmented inverse problem is formulated based on bijective mapping by introducing an additional latent random variable z .

Theory - INN - Transformations

Reversible Transformations

$$y = f(x) = x \odot \exp(s(x)) + t(x) \quad (6)$$


$$x = f^{-1}(y) = (y - t(x)) \odot \exp(-s(x)) \quad (7)$$

The design of $s(x)$ and $t(x)$ is critical. They must be constructed in a way that their outputs do not depend on all components of x , thus allowing for the computation of the inverse. Achieved through channel-wise splitting:

$$\begin{aligned} v_1 &= u_1 \exp(s_1(u_2)) + t_1(u_2) \\ v_2 &= u_2 \exp(s_2(v_1)) + t_1(v_1) \end{aligned} \quad (8)$$

s_j and t_j scaling and transformation functions modeled by neural networks. Evaluated in the forward direction, even if the block is inverted.

Method - INN - Training



```
1  for x, y in train_loader:
2      optimizer.zero_grad()
3
4      # Forward
5      output, _ = model(x_padded)
6      y_pred, z = output[:, :ndim_y], output[:, ndim_y:]
7      loss = loss_fit(y_pred, y)
8      loss += loss_latent(z)
9
10     # Backward
11     x_pred, _ = model(y, rev=True)
12     loss += loss_rev_fit(x_pred, x)
13
14     loss.backward()
15     optimizer.step()
```

1. Forward MSE
2. Latent MMD
3. Backward MSE

Method - INN - Challenges

Results & Discussion

Results - MDN - Model outputs

Model output




Figure: Model Figure: Model Figure: Model Figure: Model

3

4

5

6

Results - MDN - Performance metrics

- Evaluating performance in forward pass for all models as well as the average model.

Model	RMSE	Log MAE
Model 1	0.04744	-1.47919
Model 2	0.18287	-0.32174
Model 3	0.03565	-1.63800
Model 4	0.07076	-1.72049
Model 5	0.14245	-0.46148
Model 6	0.04430	-2.03734
Overall Average RMSE		0.07653
Overall Average Log MAE		-1.47330

Table: Performance Metrics of Models

Results - MDN - Visualizing predictions

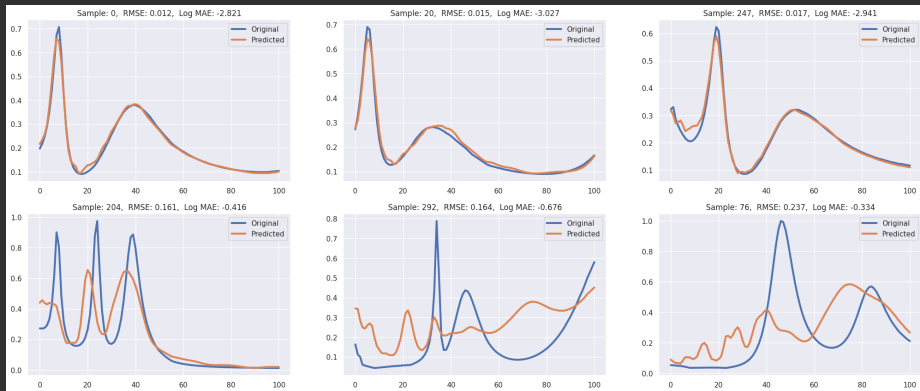


Figure: Caption

Result - DPGMM - Performance Metrics

- Evaluating the DPGMM's performance in forward and inverse analyses.

Metric	Metric Type	Value
MSE	Forward Analysis	0.013258
MSE	Inverse Analysis	0.058056
Training Loss	-	0.163675
Validation Loss	-	0.480573

Table: Performance Metrics for the DPGMM

Result - DPGMM - PDF Plot Visualizations

- Visualization of the density distribution of DPGMM's predictions.
- True Value, Mean, and Median Predictions are indicated by colored lines.




Figure: Probability Density Function (PDF) Plot of DPGMM Predictions

Results - INN - Quantitative

Metric Type	Metric	Value
Prediction	Forward MSE	3.29×10^{-4}
	Backward MSE	8.41×10^{-3}
Training	Train Loss	0.0462
	Validation Loss	0.0793
Reverse	Reverse Train Loss	0.0018
	Reverse Val Loss	0.0809

Table: Performance Metrics for Invertible Neural Network (INN). Loss on INN refers to the weighted composition loss.

Results - INN - Qualitative




Figure: Density distribution of generated outputs for sample 0 across five dimensions.

Discussion - INN

Future Work - INN

1. Active Learning Integration
2. Physics-Informed Neural Networks
3. Transfer Learning Applications

Conclusions

Conclusions

The End